

CS 486/686 Assignment 4 (64 marks)

Alice Gao

Due Date: 11:59 PM ET on Monday, December 7, 2020
with an extension without penalty
to 11:59 PM ET on Thursday, December 10, 2020

Changes

v1.1

- Corrected the total marks to 64 instead of 36.
- Corrected the formulas in Q2. Changed $\sum_{a'}$ to $\sum_{s'}$.

Academic Integrity Statement

I declare the following statements to be true:

- The work I submit here is entirely my own.
- I have not shared and will not share any of my code with anyone at any point.
- I have not posted and will not post my code on any public or private forum or website.
- I have not discussed and will not discuss the contents of this assessment with anyone at any point.
- I have not posted and will not post the contents of this assessment and its solutions on any public or private forum or website.
- I will not search for assessment solutions online.
- I am aware that misconduct related to assessments can result in significant penalties, possibly including failure in the course and suspension (this is covered in Policy 71: <https://uwaterloo.ca/secretariat/policies-procedures-guidelines/policy-71>).

Failure to accept the integrity policy will result in your assignment not being graded.

By typing or writing my full legal name below, I confirm that I have read and understood the academic integrity statement above.

Instructions

- Submit the assignment in the A4 Dropbox on LEARN. No late assignment will be accepted. This assignment is to be done individually.
- For any programming question, we will run your code in the CS-Teaching environment (linux.student.cs.uwaterloo.ca). You should make sure that (1) the language you use is available on this environment by default. (2) your code runs in this environment. We highly recommend using Python.
- Lead TAs:
 - Ethan Ward (e7ward@uwaterloo.ca)
 - Paulo Pacheco (ppacheco@uwaterloo.ca)

The TAs' office hours will be posted on MS Teams.

- Submit two files with the following names.
 - **writeup.pdf**: Include your name, email address and student ID in the writeup file. If you hand-write your solutions, make sure your handwriting is legible and take good quality pictures. You may get a mark of 0 if we cannot read your handwriting.
 - **code.zip**: Include your program, and a script to run your program. The TA will run your program using the following command in a terminal in the environment `linux.student.cs.uwaterloo.ca`.

```
bash a4.sh
```

You may get a mark of 0 if we cannot run your program.

Learning goals

Decision Networks

- Model a real-world problem as a decision network with sequential decisions.
- Given a decision network with sequential decisions, determine the optimal policy and the expected utility of the optimal policy by applying the variable elimination algorithm.

Markov Decision Process

- Trace the execution of and implement the value iteration algorithm to solve a Markov decision process.

1 Decision Network for “Monty Hall” (28 marks)

The Monty Hall Problem is stated as follows.

You are on a game show, and you are given the choice of three doors: Behind one door is a car; behind the others, goats. The host knows what’s behind each door but you don’t.

- First, you pick a door, say Door 1.
- Second, the host opens another door, say Door 3, which has a goat behind it.
- Finally, the host says to you, “Do you want to pick Door 2?”

Is it to your advantage to switch your choice?

The host always opens a door with a goat behind it, but not the door you chose first, regardless of which door you chose first. You “reserve” the right to open the door you chose first, but can change to the remaining door after the host opens the door to reveal a goat. You get as a prize the item behind the final door you choose. You prefer cars over goats (cars are worth 1 and goats are worth 0). The car is behind doors 1, 2, and 3 with probabilities p_1 , p_2 and $1 - p_1 - p_2$ respectively, and you know the values of p_1 and p_2 .

Model this problem using a decision network using the following variables.

- $\text{CarDoor} \in \{1, 2, 3\}$ is the door such that the car is behind it. This is a random variable.
- $\text{FirstChoice} \in \{1, 2, 3\}$ is the index of the door you picked first. This is a decision variable.
- $\text{HostChoice} \in \{\text{smaller}, \text{bigger}\}$ is the index of the door picked by the host. The value of this variable indicates whether the index of the door picked is the smaller or bigger one of the two doors left after you made your first choice. This is a random variable.
- $\text{SecondChoice} \in \{\text{stay}, \text{switch}\}$ indicates whether you stay with the door you picked first or switch to the remaining door not opened by the host. This is a decision variable.
- $\text{Utility} \in \{0, 1\}$ is 0 if you get a goat and 1 if you get a car.

Please complete the following tasks.

1. Complete the decision network in Figure 1 by drawing all the arcs. Show the probability table for each random variable. Show the utility table for the utility node. You can use p_1 and p_2 in your tables since you do not know their values yet.

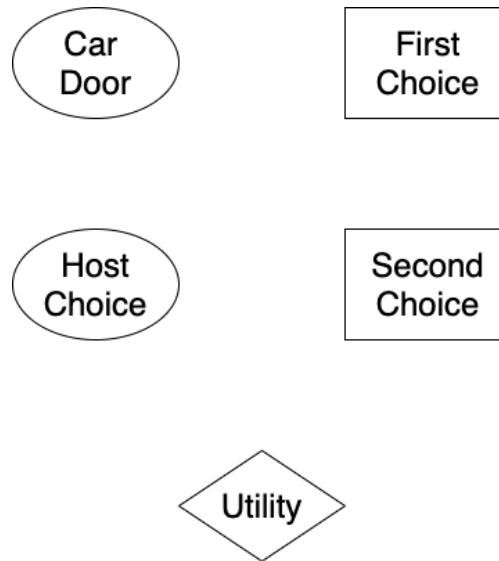


Figure 1: The Monty Hall Problem

Marking Scheme:

(10 marks)

- (4 marks) Correct parent nodes.
- (3 marks) Correct probability tables.
- (3 marks) Correct utility table.

2. Assume that $p_1 = 1/3$ and $p_2 = 1/3$.

Compute the optimal policy for the decision network by applying the variable elimination algorithm. Show all your work including all the intermediate factors created. Clearly indicate the optimal policy and the expected utility of the agent following the optimal policy.

Marking Scheme:

(9 marks)

- (4 marks) Sum out variables in the correct order.
- (2 marks) Correct optimal policy for FirstChoice.
- (2 marks) Correct optimal policy for SecondChoice.
- (1 mark) Correct value of the expected utility of the optimal policy.

3. Consider a different case where $p_1 = 0.8$ **and** $p_2 = 0.1$. The car is much more likely to be behind door 1 than to be behind doors 2 or 3.

Compute the optimal policy for the decision network by using the variable elimination algorithm. Show all your work including all the intermediate factors created. Clearly indicate the optimal policy and the expected utility of the agent following the optimal policy.

Marking Scheme:

(9 marks)

- (4 marks) Sum out variables in the correct order.
- (2 marks) Correct optimal policy for FirstChoice.
- (2 marks) Correct optimal policy for SecondChoice.
- (1 mark) Correct value of the expected utility of the optimal policy.

2 Reinforcement Learning (36 marks)

You will explore two grid worlds similar to the one discussed in lecture by running the active version of the adaptive dynamic programming algorithm.

We will test your program on two grid worlds. See their descriptions in section 2.1.

Implement the **adaptive dynamic programming** algorithm for **active reinforcement learning**. Check section 2.2 for details and tips on implementing the active ADP algorithm.

Please complete the following tasks.

When the TA runs `bash a4.sh`, the script should run **active ADP on both the lecture world and the a4 world**. The script should print out the **long-term utility values and the optimal policy for each world**.

1. For the grid world **lecture**, report the **long-term utility values** for all the states learned by the active ADP algorithm.

Because of randomness in the problem, your answers do not have to match our answers exactly. We will determine a range of acceptable values. As long as your values are within this range, you will get full marks.

Marking Scheme: (6 marks) Long-term utility values within the acceptable range.

2. For the grid world **lecture**, report the **optimal policy** given the long-term utility values.

Similar to the previous part, we will mark your answers based on whether they match multiple possible answers.

Marking Scheme: (12 marks)

- (6 marks) The TA can run your program to reproduce an optimal policy that matches a possible answer.
- (6 marks) Optimal policy is one of the possible answers.

3. For the grid world **a4**, report the **long-term utility values** for all the states learned by the active ADP algorithm.

Similar to the previous parts, we will mark your answers based on whether they fall within the acceptable range of values.

Marking Scheme: (6 marks) Long-term utility values within the acceptable range.

4. For the grid world **a4**, report the **optimal policy** given the long-term utility values. Similar to the previous parts, we will mark your answers based on whether they match multiple possible answers.

Marking Scheme: (12 marks)

- (6 marks) The TA can run your program to reproduce an optimal policy that matches a possible answer.
- (6 marks) Optimal policy matches a possible answer.

2.1 The two grid worlds

We will test your program on two grid worlds. The **lecture** world is a version of the grid world discussed in lecture. The **a4** world is created for this assignment.

The grid world lecture

The grid world discussed in lecture is given below. S_{11} is the initial state. S_{13} and S_{23} are goal states. S_{11} is a wall.

	0	1	2	3
0				
1		X		-1
2				+1

For this world, you can assume the following.

- The immediate reward of entering any non-goal state is -0.04 .
- The transition probabilities: The agent moves in the intended direction with probability 0.8, moves to the left of the intended direction with probability 0.1, and moves to the right of the intended direction with probability 0.1.
- The discount factor is 1.

The grid world a4

- Each grid world has 4 columns and 3 rows.
- Each world has at least one goal state. Entering any goal state causes the agent to exit the world.
- The reward of entering a goal state is +1 or -1.
- The maximum immediate reward of entering a goal state is 2.
- In a given world, there is a fixed immediate reward of entering any non-goal state. The immediate reward for entering every non-goal state is the same.
- In each state, there are four available actions: up, down, left, and right.
- If the agent moves in a direction and hits a wall, the agent will stay in the same square.
- By taking an action, it is possible for the agent to reach any of the four neighbouring squares (if there is a wall on any side, the neighbouring square is the current square the agent is in.).

The information for both grid worlds is provided through `run_world.pyc`. The `run_world.pyc` was produced with Python 3.7.6. The file includes the following functions. The world string can be `lecture` or `a4`.

- `get_gamma(world)`: Returns the discount factor given the world string.
- `get_reward(world)`: Returns the immediate reward of entering any non-goal state given the world string.
- `read_grid(world)`: Returns a numpy array representing the grid given the world string.

In the grid, the meanings of the symbols are as follows.

- S is the initial state.
 - * is any other non-goal state.
 - X is a wall.
 - If the state is 1 or -1, it is a goal state with the number as the reward of entering the goal state.
- `make_move(grid, curr_state, dir_intended, world)`: Given the grid, the current state, the intended direction and the world string, make a move based on the transition probabilities and return the next state.

A state is encoded as a tuple (i,j) where i is the row index ($0 \leq i \leq 2$) and j is the column index ($0 \leq j \leq 3$).

The intended direction is 0 for up, 1 for right, 2 for down, and 3 for left.

In addition, `run_world.pyc` has a few other helper functions that you can use:

- `get_next_states(grid, curr_state)`: Given a grid and the current state, return the next states for all actions. The returned array contains the next states for the actions up, right, down, and left, in this order. For example, the third element of the array is the next state if the agent ends up going down from the current state.
- `is_goal(grid, state)`: Returns true if the given state is a goal state and false otherwise. The `grid` parameter needs to be the one returned by `read_grid`. This function checks whether the given state (i, j) is 1 or -1.
- `is_wall(grid, state)`: Returns true if the given state is a wall and false otherwise. The `grid` parameter needs to be the one returned by `read_grid`. This function checks whether the given state (i, j) is 1 or -1.
- `not_goal_and_wall(grid, state)`: Returns true if the given state is not a goal state and nor a wall.
- `pretty_print_policy(grid, policy)`: Prints a policy in a readable format where the actions are printed as up, right, down, and left. The `grid` parameter needs to be the one returned by `read_grid`. The `policy` is a similar grid except that each element is 0, 1, 2, 3. 0 is up, 1 is right, 2 is down, and 3 is left.

2.2 Tips for implementing active ADP

The main idea of active ADP are the following.

- For each observed transition, update our estimates of the immediate reward values and the transition probabilities.
- Solve for the long-term expected utilities of the states (exactly or iteratively) by using the Bellman equations given our current estimates of the immediate reward values and the transition probabilities.
- Solve for the optimal policy given the current estimates of the long-term expected utilities of the states.
- Given the optimal policy and the transition probabilities, make a move by determine the next state.
- Repeat the steps until the utility values converge.

We recommend that you keep track of the following information.

- $N(s, a)$: the number of times that we have visited a given state-action pair.
- $N(s, a, s')$: the number of times we have reached state s' by taking action a in state s .
- $R(s)$: The immediate reward of entering state s .
- $U(s)$: The long-term expected utility of entering state s and following the optimal policy thereafter.

See a more detailed description of the active ADP algorithm below.

1. Given the grid, the current utility estimates, the current state, and the counts $N(s, a)$ and $N(s, a, s')$, determine the best action as follows.

$$\arg \max_a f \left(\sum_{s'} P(s'|s, a) U(s'), N(s, a) \right) \quad (1)$$

$$f(u, n) = \begin{cases} R^+, & \text{if } n < N_e \\ u, & \text{otherwise.} \end{cases} \quad (2)$$

The meanings of some expressions are explained below.

- $U(s)$ is the long-term total expected reward of entering state s and following the optimal policy thereafter.
- $R(s)$ is the immediate reward of entering state s .
- $P(s'|s, a)$ is the probability of transitioning to state s' if the agent executes action a in state s . The transition probability $P(s'|s, a)$ can be estimated as $N(s, a, s')/N(s, a)$.
- $N(s, a)$ is the number of times that the agent has executed action a in state s .
- R^+ is an optimistic estimate of the best possible reward obtainable in any state. You can use $R^+ = 2$ since the maximum immediate reward of entering any state is 2.
- N_e is a fixed parameter. This update equation ensures that the agent will try each state-action pair at least N_e times. Use $N_e = 30$.

The utility estimates $f(\sum_{s'} P(s'|s, a) U(s'), N(s, a))$ are called the optimistic estimates of the long-term utility values. If we have not tried a state-action pair (s, a) for at least N_e times, the algorithm will assume that the expected utility of the state-action pair is R^+ , which is the maximum possible reward obtainable in any state. This ensures that we will perform enough exploration by trying each state-action pair at least N_e times.

2. Make a move and determine the next state based on the current state, the best action, and the current transition probabilities. You can do this by calling the `make_move` function in `run_world`.
3. Update $N(s, a)$ and $N(s, a, s')$.
4. Update the utility values $U(s)$ by performing value iteration updates until the values converge.

The value iteration updates are as follows. Note that we are still using the optimistic utility estimates.

$$U(s) \leftarrow R(s) + \gamma \max_a f \left(\sum_{s'} P(s'|s, a) U(s'), N(s, a) \right) \quad (3)$$

$$f(u, n) = \begin{cases} R^+, & \text{if } n < N_e \\ u, & \text{otherwise.} \end{cases} \quad (4)$$

You might want to implement the following helper functions.

- `explore(u, n)`: Determine the value of the exploration function f given u and n .
- `get_prob(n_sa, n_sas, curr_state, dir_intended, next_state)`: Determine the transition probability based on counts.
- `exp_utils(grid, utils, curr_state, n_sa, n_sas)`: Calculate the expected utilities $\sum_{s'} P(s'|s, a) U(s')$
- `optimistic_exp_utils(grid, utils, curr_state, n_sa, n_sas)`: Return the optimistic expected utilities $f \left(\sum_{s'} P(s'|s, a) U(s'), N(s, a) \right)$
- `update_utils(grid, utils, n_sa, n_sas, gamma)`: Perform value iteration updates to the long-term expected utility estimates until the estimates converge.
- `utils_to_policy(grid, utils, n_sa, n_sas)`: Determine the optimal policy given the current long-term utility value for each state.